

Friday, 17 April 2026 — Day Plan

Mood: Shit day, shit weather (81°F, hot and muggy — Vivek's nightmare range). Capped goals. Not a "big day." Wake up late, minimise outdoor time, deliver the 4:30 PM meeting with Suresh cleanly.

Key obligations

- Annoying call with one guy
- 4:30 PM meeting with Suresh
- Possible dinner with friend

Step-by-step plan (in order)

Step 1 — Wake up late → coffee shop → review highlighted papers

- Wake around 11 AM.
- Walk to coffee shop.
- Review the six Cat-1/Cat-2 annotated PDFs produced over the last two days. ~30 minutes of browsing to refresh what's there.

Step 2 — Desk work at the lab → NotebookLM video summaries

- Walk to lab.
- Watch the NotebookLM video summaries of each paper. ~4 minutes per paper × ~5 papers = 20–30 minutes of video total. One sitting.

Step 3 — NotebookLM audio summary (if not done tonight)

- **Only if you didn't do it Thursday night before sleep.** If the audio played while falling asleep last night, skip this.
- Otherwise: walking task. Put it on while walking around campus in the afternoon.
- **Purpose:** ambient exposure to the narrative Suresh has of these papers, so you know what to engage with or correct in the meeting. The paper narrative is known to be misleading / hallucinated — do NOT engage critically, just hear it.

Step 4 — Annoying call with the guy

- Irritating. Unavoidable. Just get it done.

Step 5 — 4:30 PM meeting with Suresh

- The main event of the day.

- Present the paper arc: six papers read with Cat-1/Cat-2 skills, annotated PDFs, color-coded highlights, main points.
- Focus on the YOUR OWN highlighted papers, not the NotebookLM summaries (which are misleading).
- **Hold back:** the instrumentation angle. Save for next meeting.

Step 6 — Dinner with friend (if it happens)

- After the meeting. Show up, be decent company, come home.

Why this feels draining rather than restful — the honest frame:

The earlier instinct that "casual human contact is social rest" is **WRONG** for this specific cognitive profile. For someone operating at high cognitive frequencies, low-frequency interaction isn't rest — it's *active labor without intellectual return*.

The viscous-fluid feeling is neurochemically real: when dopamine and norepinephrine calibrate to high-intensity thinking, dropping to lower intensity creates a mismatch that feels physically wrong, not just boring. Like an athlete in peak training feeling lethargic when forced to rest.

And the dinner itself isn't neutral time — it's cognitive labor: holding back thoughts, slowing speech, avoiding topics, managing reactions, staying patient through logical errors. Output without input.

This friendship is a **consumptive** friendship, not a **resonant** one (the sociology-of-friendship distinction). Consumptive: you give, they take. Resonant: both give, both receive. Consumptive friendships drain because the cognitive gradient only flows one way. The friend doesn't feel the drain because he doesn't have the same intensity to manage.

Specific friction point: when the friend calls Claude "a servant bound to obey" (or variants), it hits three nerves simultaneously — reveals he doesn't understand the collaboration model, implicitly dismisses the infrastructure Vivek has built, and puts Vivek in a no-win position (explain = exhausting; correct = condescending; let it stand = complicit). None of the three options advances the interaction.

Mitigations for Friday's dinner:

- Shorten if possible — 1 to 1.5 hours, not 3.
- Venues with parallel activity — eating while watching something — so interaction isn't continuous.
- Avoid AI as a topic. It's a tar pit.
- Reframe expectation: this is transactional maintenance of a legacy relationship, not a nourishing event.
- Get in, be decent company, get out, don't let it bleed into the evening.
- Don't resent the friend — he's not operating at a lower frequency maliciously, he just *is* at a different frequency. The mistake is expecting this specific dinner to feed intellectual life; the cure is not changing the friend, it's adjusting expectations.

The consumer-creator axis (separate from the frequency axis):

The friction has a second, independent dimension. The friend's default orientation to reality is *consumptive* ("is there a tutorial?" "can I watch this guy talk to AI?" "I'll look at it eventually when everything comes out"). Vivek's default orientation is *creative* (go talk to AI yourself, build from principles, make the thing that others will eventually consume).

Why this matters uniquely NOW: the current AI moment rewards creators disproportionately because there's nothing mature to consume yet. The field is too new for tutorials. People thriving right now are building from principles; people waiting are stranded. Waiting for "it to come out" means missing the window in which being a creator has outsized leverage.

Why consumer-to-creator interaction is energetically asymmetric:

- You can tell a consumer "go create," but you can't *make* them create. The mode isn't contagious in that direction.
- Their consumptive suggestions ("watch this play," "look at this guy on YouTube") require mode-switching labor on your part even to engage with.
- Over time the gap *widens*, not closes. Consumers stay consumers. Creators keep creating. The two modes diverge.

Related frameworks: Paul Graham on makers vs managers (adjacent). Taleb's doers-vs-talkers. Austin Kleon's creator-disposition essays. It's a real and recognized split in how people orient to knowledge and experience.

What this means for the friendship: the gap won't heal itself by time. Expect the friendship to remain as it is, or slowly drift further apart as the creator-consumer delta grows. That's okay. Not every friendship needs to scale with you — some are historical, some are consumptive-maintenance, some are resonant. Know which one this is and don't expect it to become something else.

Held for next meeting / later

Instrumentation discussion

- The real payoff from the Chen + Ergodic papers is **instrumentation**, not methodology.
- Chen's custom carbon-fibre sensor → reference [35], Hu et al. 2021. Still needs to be read separately.
- Ergodic's commercial F/T sensor (Bota SensOne) → opened the door to the Bota vs GelSight vs Tekscan FlexiForce sensor brainstorm.
- Investigation is ongoing; conclusions are still forming. Need a few more papers read and more brainstorming before bringing this to Suresh.
- **Next-meeting agenda item.**

Project Genie / world-model prototyping — OPEN CASE, ongoing exploration

- Explored in tonight's voice session. Five prompt iterations on a hemisphere + handheld sensor scenario. Culminated in the chicken iPhone.
- **Do NOT conclude from one session.** Single-session judgment on a brand-new tool with its own prompt dialect is over-fitting. The chicken iPhone is a data point, not a verdict.
- **Why this is worth sustained investment:**
 - Suresh paid \$250 for the Google AI Ultra subscription. Financial commitment = exploration obligation.
 - DeepMind claims Genie 3 has *emergent physics*. Verifying or falsifying that claim is a research question, not a tool-usage question.
 - Connects to the Computational Physics 2.0 vision (to be discussed separately with Claude).
- **Methodology — physics-probe approach:**
 - Design 5–6 tiny scenarios that each isolate one physics phenomenon: 1. Object permanence (does a ball stay same size after look-away and look-back?) 2. Gravity (does a dropped object accelerate?) 3. Two-object collision (do two balls interact physically?) 4. Soft deformation (does a squishy object deform and spring back correctly?) 5. Material identity (does "steel" behave differently from "rubber" in same scenario?) 6. Continuous motion (does smooth motion look smooth?)
 - Each probe is a small Genie session. Learn what Genie EMERGES vs FAILS across categories.
 - Output: empirical characterisation of Genie 3's emergent-physics claim. This IS research, presentable to Suresh.
- **Sora's role:** validation only, not primary. For any probe where Genie looks weird, prompt Sora with the same scene to get a cinematic reference. Helps distinguish "Genie's physics model is failing" from "my prompt is wrong."
- **Avoid the A/B trap:** don't write Sora-native prompts and drop them into Genie. That evaluates Genie on Sora's axis. Start fresh with Genie-native prompting (interactive systems, consistent over-time behaviour), use Sora only as a secondary check.

Saturday payoff

- Saturday forecast: 49°F high, 34°F low, mostly clear, calm evening. Near-ideal conditions.
- Target: 40K+ steps spread across the whole day. Take advantage of the entire cool window — not just the evening.
- POWER MODE remote control for complex Claude Code work during that evening window.
- Friday is just the bridge.